





MIPS

Assembly Programming



MIPS Assembly Programming

Robert Winkler

Version 1.0.14, 2026-09-29

Table of Contents

Info	1
Dedication	2
Chapter 0: Hello World	3
Prereqs	3
System Setup	3
Handy Resources	4
Hello World	4
Building and Running	6
Conclusion	6
Exercises	7
Chapter 1: Data	8
Arrays	9
Exercises	10
Chapter 2: System Calls	12
Examples	13
Exercises	15
Chapter 3: Branches and Logic	17
Practice	18
Conclusion	27
Exercises	27
Chapter 4: Loops	28
Looping Through Arrays	30
Conclusion	35
Exercises	35
Chapter 5: Functions and the MIPS Calling Convention	37
Functions	37
The Convention	38
Conclusion	44
Exercises	44
Chapter 6: Floating Point Types	46
Floating Point Registers and Instructions	46
Practice	47
Getting Floating Point Literals	48
Branching	48
Functions	49
Conclusion	50
Exercises	50
Chapter 7: Tips and Tricks	51

Formatting	51
Misc. General Tips	51
Constants	52
Macros	53
Switch-Case Statements	54
Command Line Arguments	58
Delayed Branches and Delayed Loads	62
No Pseudoinstructions Allowed	63
Exercises	67
References and Useful Links	69
Supporters	70
Corporate	70

Info

Copyright © 2021-2026 [Robert Winkler](#)

This book is licensed under the [Creative Commons BY-NC-SA 4.0](#) which summarized means:

You are free to:

- **Share** — copy and redistribute the material in any medium or format
- **Adapt** — remix, transform, and build upon the material
- The licensor cannot revoke these freedoms as long as you follow the license terms.

Under the following terms:

- **Attribution** — You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.
- **NonCommercial** — You may not use the material for commercial purposes.
- **ShareAlike** — If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original.
- **No additional restrictions** — You may not apply legal terms or technological measures that legally restrict others from doing anything the license permits.

This book is available online in [HTML](#), [PDF](#), and [EPUB](#) formats.

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

The repo for the book, where you can get the code referenced and report any errors (submit an issue or even a pull request) is [here](#). The code is MIT licensed.


If you're interested in contacting me regarding MIPS tutoring or any other business request related to the book, you can reach me at books@robertwinkler.com.

Dedication

This book is dedicated to all its [supporters](#) and all the students I've helped with MIPS over the years who inspired me to create it.

Thank you.

An extra thank you goes to the corporate level sponsors below:



Chapter 0: Hello World

In which we lay the groundwork for the rest of the book...

Prereqs

While someone with no programming experience could probably learn MIPS from this book, it is definitely preferable to have at least some experience in a higher level imperative programming language. I say imperative, because programming in assembly is the antithesis of functional programming; everything is about state, with each line changing the state of the CPU and sometimes memory. Given that, experience in functional languages like Lisp, Scheme etc. are less helpful than experience in C/C++, Java, Python, Javascript etc.

Of all of the latter, C is the best, with C++ being a close second because at least all of C exists in C++. There are many reasons C is the best prior experience when learning assembly (any assembly, not just MIPS), including the following:

- pointers, concepts and symmetry of "address of" and "dereference" operators
- pointer/array syntax equivalence
- stack allocation as the default
- manual memory management, no garbage collector
- global data
- rough equivalence in structure of a C program and an assembly program (vs. say Java)
- pass by value

There is some overlap between those and there are probably more, but you can see that most other languages that are commonly taught as first languages are missing most, if not all of those things.

Even C++, which technically has all of them, being a superset of C, is usually taught in a way that mostly ignores all of those things. They teach C++ as if it's Java, never teaching the fundamentals. In any case this is getting into my problems with CS pedagogy of the last 20 years based on my experience as a CS major myself ('12) and as a programming tutor helping college students across the country since 2016, and I should save it for a proper essay/rant sometime.

Long story short, I use C code and C syntax to help explain and teach MIPS. I'll try to provide enough explanation regardless of past experience as best I can.

System Setup

As I tell all of my tutoring students, if you're majoring in CS or anything related I highly recommend you use Linux. It's easier in every way to do dev work on Linux vs Windows or Mac. Many assignments require it, which often necessitates using a virtual machine (which is painful, especially on laptops) and/or ssh-ing into a school Linux server, which is also less than ideal. In general, you'll have to learn how to use the Unix terminal eventually and will probably use it to some extent in your career so it also makes sense to get used to it asap.

That being said, Windows does now have WSL so you can get the full Ubuntu or Debian or Fedora etc. terminal based system on Windows without having to setup a real virtual machine (or dealing with the slowdown that would cause). I've even heard that they'll get support for Linux GUI programs soon.

MacOS on the other hand, is technically a Unix based system and you can use their terminal and install virtually any program from there using Macports or Homebrew or similar.

There are 3 commonly used MIPS simulators that I know of:

- SPIM is probably the oldest and is terminal only
- QtSpim is a GUI front end for SPIM
- MARS is a Java GUI based simulator with dozens of extra syscalls, syntactic sugar and features like graphics, memory mapped I/O, etc.

SPIM and QtSpim are in the Debian/Ubuntu repos so you can install them with the following

```
$ sudo apt install qtspim spim
```

You can probably find them in the repos for other distros too and install them similarly.

If it's not in your repos or you're on Windows (and not using WSL) or Mac you can download them, and MARS at these addresses

- [MARS](#)
- [QtSpim](#)

Handy Resources

There are a few references that you should bookmark (or download) before you get started. The first is the [MIPS Greensheet](#). It's possible you already have a physical copy of this as it's actually the tearout from the Patterson and Hennessey textbook *Computer Organization and Design MIPS Edition* that is commonly used in college courses.

The second thing is the list of [syscalls](#) from the MARS website.

I recommend you download/bookmark both and keep them open while working because you'll be referencing them often to remind yourself which instructions and syscalls you have available and how they work.

Hello World

Let's start with the classic hello world program, first in C, then in MIPS, and go over all the pieces in overview. You can copy paste these into your editor of choice (mine being neovim), or use the files in the associated repo to follow along.

```
1 #include <stdio.h>
```

```

2
3 int main()
4 {
5     printf("Hello World!\n");
6     return 0;
7 }

```

It is pretty self explanatory. You have to include `stdio.h` so you can use the function `printf` (though in the real world I'd use `puts` here), the function `main` is the start of any C/C++ program, which is a function that returns an `int`. We call `printf` to display the string "Hello World!\n" to the user and then return 0 to exit. Returning 0 indicates success and there were no errors.

You can compile and run it in a linux/unix terminal as shown below. You can substitute `clang` or another compiler for `gcc` if you want.

```

$ gcc -o hello hello.c
$ ./hello
Hello World!

```

Now, the same program in MIPS:

```

1 .data
2 hello: .asciiz "Hello World!\n"
3
4 .text
5 main:
6     li $v0, 4      # load immediate, v0 = 4 (4 is print string system call)
7     la $a0, hello  # load address of string to print into a0
8     syscall
9
10    li $v0, 10     # exit syscall
11    syscall

```

The `.data` section is where you declare global variables, which includes string literals as in this case. We'll cover them in more detail later.

The `.text` section is where any code goes. Here we declare a single label `main:`, indicating the start of our main function.

We then put the number 4 in the `$v0` register to select the print string system call. The print string system call takes one argument, the address of the string to print, in the `$a0` register. We do that on the next line. On line 8, we call the system call using the `syscall` instruction.

Finally we call the exit system call which takes no arguments and exits the program.

Again, we'll cover system calls in a later chapter. This is just an intro/overview so don't worry if some things aren't completely clear. This chapter is about getting you up and running, not really about teaching anything specific yet.

Building and Running

Now that we have our hello world MIPS program, how do we run it? Well the easiest and quickest^[1] way is of course to do it on the command line, which can be done like this for spim:

```
$ spim -file hello.s
SPIM Version 8.0 of January 8, 2010
Copyright 1990-2010, James R. Larus.
All Rights Reserved.
See the file README for a full copyright notice.
Loaded: /usr/lib/spim/exceptions.s
Hello World!
```

or this for MARS:

```
$ java -jar ~/Mars4_5.jar hello.s
MARS 4.5 Copyright 2003-2014 Pete Sanderson and Kenneth Vollmar

Hello World!
```

The name of your MARS jar file may be different^[2], so be sure to use the correct name and path. For myself, I keep the jar file in my home directory so I can use tilde to access it no matter where I am. You can also copy it into your working directory (ie wherever you have your source code) so you don't have to specify a path at all. There are lots of useful command line options that you can use^[3], some of which we'll touch on later.

Running the jar directly on the command line works even in the DOS command line though I've never done it and it's probably not worth it.

Alternatively, you can start up MARS or QtSpim like a normal GUI application and then load your source file. MARS requires you to hit "assemble" and then "run". Whereas with QtSpim you only have to hit "run".

QtSpim does let you start and load the file in one step from the command line

```
$ qtspim hello.s
```

but there is no way to simply run it with out starting the GUI, which makes sense since the whole point is to be a GUI wrapper around spim.

Conclusion

Well, there you have it, you have written and run your first MIPS program. Another few chapters and you will have no trouble with almost anything you would want to do in MIPS, whether for a class, or on your own for fun.

Exercises

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

1. Modify the hello world program to print something different, perhaps your name.
2. Run it with both SPIM and MARS.

[1] Starting up the MARS GUI (an old Java app) is often annoyingly slow

[2] Some schools/professors have their own versions with extra features and other improvements over the old version available on the MARS website

[3] <https://dpetersanderson.github.io/Help/MarsHelpCommand.html>

Chapter 1: Data

In MIPS, you can declare global variables in the `.data` section.

At a minimum, this is where you would declare/define any literal strings your program will be printing, since virtually every program has at least 1 or 2 of those.

When declaring something in the `.data` section, the format is

```
variable_name: .directive value(s)
```

where whitespace between the 3 is arbitrary. The possible directives are listed in the following table:

Table 1. MIPS data types

Directive	Size (bytes)	C equivalent
<code>.byte</code>	1	char
<code>.half</code>	2	short
<code>.word</code>	4	int, all pointer types
<code>.float</code>	4	float
<code>.double</code>	8	double
<code>.ascii</code>	NA	char str[5] = "hello"; (no '\0')
<code>.asciiz</code>	NA	char str[] = "hello"; (includes the '\0')
<code>.space</code>	NA	typeless, uninitialized space, can be used for any type/array

As you can see it's pretty straightforward, but there are a few more details about actually using them so let's move onto some examples.

Say you wanted to convert the following simple program to MIPS:

```
1 #include <stdio.h>
2
3 int main()
4 {
5     char name[30];
6     int age;
7     printf("What's your name and age?\n");
8     scanf("%s %d", name, &age);
9     printf("Hello %s, nice to meet you!\n", name);
10    return 0;
11 }
```

The first thing you have to remember when converting from a higher level language to assembly (any assembly), is that what matters is whether it is functionally the same, not whether everything

is done in exactly the same way.^[1] In this instance, that means realizing that your literal strings and your local variables `name` and `age` become globals in MIPS.

```
1 .data
2 age:      .word 0 # can be initialized to anything
3
4 ask_name:  .asciiz "What's your name and age?\n"
5 hello_space: .asciiz "Hello "
6 nice_meet: .asciiz ", nice to meet you!\n"
7
8 name:      .space 30
9
10 .text
11
12 # main goes here
```

As you can see in the example, we extract all the string literals and the character array `name` and int `age` and declare them as MIPS globals. One thing to note is the second `printf`. Because it prints a variable, `name`, using the conversion specifier, we break the literal into pieces around that. Since there is no built-in `printf` function in MIPS, you have to handle printing variables yourself with the appropriate system calls.

Arrays

Obviously strings are special cases that can be handled with `.ascii` or `.asciiz` for literals, but for other types or user input strings how do we do it?

The first way, which was demonstrated in the snippet above is to use `.space` to declare an array of the necessary byte size. Keep in mind that the size is specified in bytes not elements, so it only matches for character arrays. For arrays of ints/words, floats, doubles etc. you'd have to multiply by the `sizeof(type)`.

"But, `.space` only lets you declare uninitialized arrays, how do I do initialized ones?"

Actually, it appears `.space` initializes everything to 0 similar to global/static data in C and C++, though I can't find that documented anywhere.

Aside from that, there are two ways depending on whether you want to initialize every element to the same value or not.

For different values, the syntax is an extension of declaring a single variable of that type. You specify all the values, comma separated. This actually gives you another way to declare a string or a character array, though I can't really think of a reason you'd want to. You could declare a `.byte` array and list all the characters individually.

However, if you want an array with all elements initialized to the same value there is a more convenient option. After the type you put the value you want, a colon, and then the number of elements. So `a: .word 123 : 10` would declare a 10 integer array with all elements set to 123. This works for all types in MARS, but SPIM does not support this syntax for floats and doubles. Since

you'll probably not have to deal with floating point types at all, let alone arrays of them initialized to a particular value, this isn't a huge loss.

Given what we just covered, this:

```
1 int a[20];
2 double b[20];
3 int c[10] = { 9,8,7,6,5,4,3,2,1,0 };
4 int d[5] = { 42, 42, 42, 42, 42 };
5 char e[3] = { 'a', 'b', 'c' };
```

becomes

```
1 .data
2 a:      .space 80
3 b:      .space 160
4 c:      .word 9,8,7,6,5,4,3,2,1,0
5 d:      .word 42 : 5
6 e:      .byte 'a', 'b', 'c'
```

For more examples of array declarations, see [array_decls.s](#). You don't have to understand the rest of the code, just that it prints out each of the arrays.

Exercises

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

1. Create a MIPS data section that declares variables equivalent to the following. This will not be a runnable program without a main.

```
1 float a;
2 float b = 2.71;
3 int myarray[10] = { 9, 8, 7, 6, 5, 4, 3, 2, 1 };
4 short array2[10];
5
6 char mips_str[] = "MIPS assembly is awesome!";
```

2. How would you declare an array of 500 points? The point structure is tightly packed and defined like this:

```
1 struct point {
2     float x;
3     float y;
4     float z;
```

```
5 };
```

[1] Obviously compilers have to follow stricter rules, but for the purposes of learning and actually using assembly directly, there's no reason to make your life harder than necessary.

Chapter 2: System Calls

We mentioned system calls (aka syscalls from now on) in chapter 0 when we were going over our "Hello World" program, but what exactly are they?

Essentially, they are the built in functions of an operating system; in this case, the simple operating system of the MIPS simulators. They provide access to all the fundamental features, like input and output to/from both the console and files, allocating memory, and exiting. That covers all the 17 syscalls supported by spim, but MARS supports many more, for things ranging from playing MIDI sounds, to getting a random number, to creating GUI dialogs.^[1]

NOTE

Except for the MARS or SPIM specific chapters/sections, I'll be sticking to code compatible with both throughout this book, meaning we only use the first 17 syscalls, and don't get to use some of the syntactic sugar available in MARS, or any SPIM specific features either.

Table 2. SPIM supported syscalls

Name	\$v0	Arguments	Result
print integer	1	\$a0 = integer to print	
print float	2	\$f12 = float to print	
print double	3	\$f12 = double to print	
print string	4	\$a0 = address of string	
read integer	5		\$v0 = integer read
read float	6		\$f0 = float read
read double	7		\$f0 = double read
read string	8	\$a0 = address of input buffer \$a1 = buffer size	works like C's fgets
sbrk	9	\$a0 = size in bytes to allocate	\$v0 = address of allocated memory (sbrk is basically malloc but there is no free)
exit	10		program terminates
print character	11	\$a0 = character to print (ascii value)	
read character	12		\$v0 = character read
open file	13	\$a0 = address of filename \$a1 = flags \$a2 = mode	\$v0 = file descriptor (negative if error)
read from file	14	\$a0 = file descriptor \$a1 = address of input buffer \$a2 = max characters to read	\$v0 = number of characters read, 0 for end-of-file, negative for error

Name	\$v0	Arguments	Result
write to file	15	\$a0 = file descriptor \$a1 = address of output buffer \$a2 = number of characters to write	\$v0 = number of characters written, negative for error
close file	16	\$a0 = file descriptor	
exit2	17	\$a0 = termination result	program terminates, returning number in \$a0 (only meaningful when run in the terminal, ignored in GUI)

As you can see, they really only cover the basics. You can read or write the different types, do file I/O using calls identical to POSIX functions (open, read, write, close; see man pages), allocate memory, and exit. Even so, they're sufficient to build anything you want.

So, what does that table mean? How do these actually work?

The process is:

1. Put the number for the syscall you want in `$v0`
2. Fill in the appropriate arguments, if any
3. Execute the syscall with `syscall`

```
1  li    $v0, 1    # 1 is print integer
2  li    $a0, 42   # takes 1 arg in a0, the number to print
3  syscall        # actually execute syscall
```

You can think of the above as `print_integer(42);`. Let's look at an actual program that uses a few more syscalls next.

Examples

```
1  #include <stdio.h>
2
3  int main()
4  {
5      int age;
6      int height;
7      char name[50];
8      printf("What's your name? ");
9      fgets(name, 50, stdin);
10
11     printf("Hello %s", name);
12
13     printf("How old are you? ");
14     scanf("%d", &age);
15
```

```

16     printf("Enter your height in inches: ");
17     scanf("%d", &height);
18
19     printf("Your age + height = %d\n", age + height);
20
21     return 0;
22 }

```

I'm using `fgets()` instead of `scanf("%s", name)` because `fgets` works the same as the `read string syscall (8)`.

```

1 .data
2
3 name:      .space 50
4
5 nameprompt: .asciiz "What's your name? "
6 hello_space: .asciiz "Hello "
7 how_old:    .asciiz "How old are you? "
8 ask_height: .asciiz "Enter your height in inches: "
9 ageplusheight: .asciiz "Your age + height = "
10
11
12 .text
13 main:
14     li    $v0, 4      # print string system call
15     la    $a0, nameprompt # load address of string to print into a0
16     syscall
17
18     li    $v0, 8      # read string
19     la    $a0, name
20     li    $a1, 50
21     syscall
22
23     li    $v0, 4
24     la    $a0, hello_space
25     syscall
26
27     la    $a0, name # note 4 is still in $v0
28     syscall
29
30     # don't print a newline here because
31     # one will be part of name
32
33     li    $v0, 4
34     la    $a0, how_old
35     syscall
36
37     li    $v0, 5      # read integer
38     syscall
39     move  $t0, $v0    # save age in t0

```

```

40
41     li    $v0, 4
42     la    $a0, ask_height
43     syscall
44
45     li    $v0, 5    # read integer
46     syscall
47     add    $t0, $t0, $v0 # t0 += height
48
49
50     li    $v0, 4
51     la    $a0, ageplusheight
52     syscall
53
54     li    $v0, 1    # print int
55     move   $a0, $t0  # a0 = age + height
56     syscall
57
58     # print newline
59     li    $v0, 11    # print char
60     li    $a0, 10    # ascii value of '\n'
61     syscall
62
63
64     li    $v0, 10    # exit syscall
65     syscall

```

There are a few things to note from the example.

We don't declare global variables for age or height. We could, but there's no reason to since we need them in registers to perform the addition anyway. Instead, we copy/save age to `$t0` so we can use `$v0` for 2 more syscalls, then add height to `$t0`.

This is generally how it works. Use registers for local variables unless required to do otherwise. We'll cover [more about](#) register use when we cover the MIPS calling convention.

Another thing is when we print their name, we don't put 4 in `$v0` again because it is still/already 4 from the lines above. Unless the syscall says it writes to `$v0` you can assume it is unmodified.

Lastly, many people will declare a string `"\n"` and use `print string` to print a newline, but it's easier to use the `print char` syscall as we do right before exiting.

Exercises

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

1. Convert the following C code to MIPS

```

1 #include <stdio.h>
2
3 int main()
4 {
5     float price;
6     double golden = 1.618;
7     int ret;
8
9     printf("Enter what the price of gas was last time you filled up: ");
10    scanf("%f", &price);
11
12    printf("%f is too expensive!\n", price);
13
14    printf("The golden ratio is roughly %f\n", golden);
15
16
17    printf("Enter an integer for the program to return: ");
18    scanf("%d", &ret);
19    return ret;
20 }

```

2. Write a program that asks the user for their name, reads it in, and then prints "Hello [user's name]!"

[1] <https://dpetersanderson.github.io/Help/SyscallHelp.html>

Chapter 3: Branches and Logic

We can't go much further in our MIPS programming journey without covering branching. Almost every non-trivial program requires some logic, even if it's only a few **if** or **if-else** statements. In other words, almost every program requires branching, a way to do a instead of b, or to do a only if certain conditions are met.

You already know how to do this in higher level languages, the aforementioned **if** statement. In assembly it's more complicated. Your only tool is the ability to jump to a label on another line based on the result of various comparisons. The relevant instructions are listed in the following table:

Table 3. MIPS branching related instructions (and pseudoinstructions)

Name	Opcode	Format	Operation
Branch On Equal	beq	beq rs, rt, label	if (rs == rt) goto label
Branch On Not Equal	bne	bne rs, rt, label	if (rs != rt) goto label
Branch Less Than	blt	blt rs, rt, label	if (rs < rt) goto label
Branch Greater Than	bgt	bgt rs, rt, label	if (rs > rt) goto label
Branch Less Than Or Equal	ble	ble rs, rt, label	if (rs <= rt) goto label
Branch Greater Than Or Equal	bge	bge rs, rt, label	if (rs >= rt) goto label
Set Less Than	slt	slt rd, rs, rt	rd = (rs < rt) ? 1 : 0
Set Less Than Immediate	slti	slt rd, rs, imm	rd = (rs < imm) ? 1 : 0
Set Less Than Immediate Unsigned	sltiu	slt rd, rs, imm	rd = (rs < imm) ? 1 : 0
Set Less Than Unsigned	sltu	sltu rd, rs, rt	rd = (rs < rt) ? 1 : 0

You can see the same information and more (like which ones are pseudoinstructions) on the MIPS greensheet.^[1]

There are additional pseudoinstructions in the form of beq/bne/blt/bgt/ble/bge + 'z' which are syntactic sugar to compare a register against 0, ie the 0 register.

So the following:

```
beq    $t0, $0, label
bne    $t1, $0, label
blt    $t2, $0, label
```

would be equivalent to:

```
beqz   $t0, label
bnez   $t1, label
bltz   $t2, label
```

Note `$0` is the same as `zero` and is the hard coded 0 register. I'll cover registers in [more detail](#) in the chapter on functions and the calling conventions.

One final thing is that labels have the same naming requirements as C variables and functions. They must start with a letter or underscore and the rest can be letters, underscores, or digits.

Practice

The rest of this chapter will be going over many examples, looking at snippets of code in C and translating them to MIPS.

Basics

Let's start with the most basic if statement. The code in and after the if statement is arbitrary.

```
1  if (a > 0) {
2      a++;
3  }
4  a *= 2;
```

Now in MIPS, let's assume that `a` is in `$t0`. The translation would look like this:

```
1  ble  $t0, $0, less_eq_0    # if (a <= 0) goto less_eq_0
2  addi $t0, $t0, 1           # a++
3  less_eq_0:
4  sll  $t0, $t0, 1           # a *= 2 (shifting left by n is multiplying by 2^n)
```

There are a few things to note in this example. The first is that in assembly we test for the opposite of what was in the if statement. This will always be the case when jumping forward because (if we want to keep the same order of code) we can only jump *over* a block of code, whereas in C we fall into the block if the condition is true. In the process of mentally compiling a bit of C to assembly, it can be helpful to change to jump based logic first. For example the previous C would become:

```
1  if (a <= 0)
2      goto less_eq_0;
3  a++;
4  less_eq_0:
5  a *= 2;
```

This is obviously still valid C but matches the branching behavior of assembly exactly. You can see I put comments for the equivalent C code in my assembly; it helps with readability to comment every line or group of lines that way.

The second thing to notice is how we handled the multiplication. This has nothing to do with branching but is something we'll touch on multiple times throughout the book. Your job when acting as a human compiler is to match the *behavior*. You are under no obligation to match the

structure or operations of the higher level code exactly (unless your professor stupidly forces you to).

Given that, it is in your best interest to change and rearrange things in order to simplify the assembly as much as possible to make your life easier. Generally speaking, this also tends to result in more performant code, since using fewer instructions and fewer branches (the most common outcomes) saves execution time.

In this case, the standard `mult` instruction (from the green sheet) would have required 3 instructions, and even the `mul` instruction (that does seem to be supported everywhere but is not on the green sheet) would take 2:

```
1    li    $t1, 2
2    mult  $t0, $t1
3    mflo  $t0          # a *= 2
4
5    # or
6
7    li    $t1, 2
8    mul   $t0, $t0, $t1 # a *= 2
```

This is why, when multiplying or dividing by a constant power of 2 it's common practice to use `sll` or `sra`. This is true in all assembly languages because multiplication and division are relatively costly operations so using shifts when you can saves performance even if you didn't actually save instructions.

Ok, let's look at an `if-else` example. Again, the actual code is arbitrary and we're assuming a and b are in `$t0` and `$t1` respectively

```
1    if (a > 0) {
2        b = 100;
3    } else {
4        b -= 50;
5    }
```

You could do it something like these two ways

```
1    bgt    $t0, $0, greater_0    # if (a > 0) goto greater_0
2    addi   $t1, $t1, -50         # b -= 50
3    j      less_eq_0
4 greater_0:
5    li     $t1, 100              # b = 100
6 less_eq_0:
7
8    # or
9
10   ble    $t0, $0, less_eq_0    # if (a <= 0) goto less_eq_0
11   li     $t1, 100              # b = 100
```

```

12     j      greater_0
13 less_eq_0:
14     addi   $t1, $t1, -50      # b -= 50
15 greater_0:

```

You can see how the first swaps the order of the actual code which keeps the actual conditions the same as in C, while the second does what we discussed before and inverts the condition in order keep the blocks in the same order. In both cases, an extra unconditional branch and label are necessary so we don't fall through the else case. This is inefficient and wasteful, not to mention complicates the code unnecessarily. Remember how our job is to match the behavior, not the exact structure? Imagine how we could rewrite it in C to simplify the logic:

```

1     b -= 50;
2     if (a > 0) {
3         b = 100;
4     }

```

which becomes

```

1     addi   $t1, $t1, -50      # b -= 50;
2     ble    $t0, $0, less_eq_0 # if (a <= 0) goto less_eq_0
3     li     $t1, 100          # b = 100
4 less_eq_0:

```

That is a simple example of rearranging code to make your life easier. In this case, we are taking advantage of what the code is doing to make a default path or default case. Obviously, because of the nature of the code subtracting 50 has to be the default since setting b to 100 overwrites the original value which we'd need if we were supposed to subtract 50 instead. In cases where you can't avoid destructive changes (like where the condition and the code are using/modifying the same variable), you can use a temporary variable; i.e. copy the value into a spare register. You still save yourself an unnecessary jump and label.

Compound Conditions

These first 2 examples have been based on simple conditions, but what if you have compound conditions? How does that work with branch operations that only test a single condition? As you might expect, you have to break things down to match the logic using the operations you have.

Let's look at **and** first. Variables a, b, and c are in t0, t1, and t2.

```

1     if (a > 10 && a < b) {
2         c += 20;
3     }
4     b &= 0xFF;

```

So what's our first step? Like previous examples, we need to test for the opposite when we switch to

assembly, so we need the equivalent of

```
1  if (!(a > 10 && a < b))
2      goto no_add20;
3  c += 20;
4 no_add20:
5  b &= 0xFF;
```

That didn't help us much because we still don't know how to handle that compound condition. In fact we've made it more complicated. If only there were a way to convert it to **or** instead of **and**. Why would we want that? Because, while both **and** and **or** in C allow for short circuit evaluation (where the result of the whole expression is known early and the rest of expression is not evaluated), with **or**, it short circuits on success while **and** short circuits on failure. What does that mean? It means that with **or**, the whole expression is true the second a single true term is found, while with **and** the whole expression is false the second a single false term is found.

Let's look at the following code to demonstrate:

```
1  if (a || b || c) {
2      something;
3  }
4
5  // What does this actually look like if we rewrote it to show what it's
6  // actually doing with short circuit evaluation?
7
8  if (a) goto do_something;
9  if (b) goto do_something;
10 if (c) goto do_something;
11 goto dont_do_something;
12
13 do_something:
14     something;
15
16 dont_do_something:
17
18     // You can see how the first success is all you need
19     // Compare that with and below:
20
21     if (a && b && c) {
22         something;
23     }
24
25     if (a) {
26         if (b) {
27             if (c) {
28                 something;
29             }
30         }
31     }
```

```

32 // which in jump form is
33
34 if (a)
35     goto a_true;
36 goto failure;
37 a_true:
38 if (b)
39     goto b_true;
40 goto failure;
41
42 b_true:
43 if (c)
44     goto c_true;
45 goto failure;
46
47 c_true:
48     something;
49 failure:
50
51 // Man that's ugly, overcomplicated, and hard to read
52 // But what if we did this instead:
53
54 if (!a) goto dont_do_something;
55 if (!b) goto dont_do_something;
56 if (!c) goto dont_do_something;
57
58     something;
59
60 dont_do_something:
61
62 // Clearly you need all successes for and. In other words
63 // to do and directly, you need state, knowledge of past
64 // successes. But what about that second translation of and?
65 // It looks a lot like or?

```

You're exactly right. That final translation of **and** is exactly like **or**.

It takes advantage of De Morgan's laws.^[2] For those of you who haven't taken a Digital Logic course (or have forgotten), De Morgan's laws are 2 equivalencies, a way to change an **or** to an **and**, and vice versa.

They are (in C notation):

`!(A || B) == !A && !B`

`!(A && B) == !A || !B`

Essentially you can think of it as splitting the not across the terms and changing the logical operation. The law works for arbitrary numbers of terms, not just 2:

`(A && B && C)`

```
is really
((A && B) && C)
so when you apply De Morgan's Law recursively you get:
!((A && B) && C) == !(A && B) || !C == !A || !B || !C
```

Let's apply the law to our current compound **and** example. Of course the negation of greater or less than comparisons means covering the rest of the number line so it becomes:

```
1  if (a <= 10 || a >= b))
2      goto no_add20;
3  c += 20;
4 no_add20:
5  b &= 0xFF;
```

which turns into:

```
1  li    $t9, 10
2  ble   $t0, $t9, no_add20    # if (a <= 10) goto no_add20
3  bge   $t0, $t1, no_add20    # if (a >= b) goto no_add20
4
5  addi   $t2, $t2, 20         # c += 20
6 no_add20:
7  andi   $t1, $t1, 0xFF      # b &= 0xFF
```

See how that works? **Or**'s do not need to remember state. Just the fact that you reached a line in a multi-term **or** expression means the previous checks were false, otherwise you'd have jumped. If you tried to emulate the same thing with an **and**, as you saw in the larger snippet above, you'd need a bunch of extra labels and jumps for each term.

What about mixed compound statements?

```
1  if (a > 10 || c > 100 && b >= c)
2      printf("true\n");
3
4  b |= 0xAA;
```

Well, the first thing to remember is that **&&** has a higher priority than **||**, which is why most compilers these days will give a warning for the above code about putting parenthesis around the **&&** expression to show you meant it (even though it's completely legal as is).

So with that in mind, let's change it to jump format to better see what we need to do. While we're at it, let's apply De Morgan's law to the **&&**.

```
1  if (a > 10)
2      goto do_true;
3  if (c <= 100)
```

```

4     goto done_if;
5     if (b < c)
6         goto done_if;
7 do_true:
8     printf("true\n");
9
10 done_if:
11     b |= 0xAA;

```

This one is trickier because we don't flip the initial expression like normal. Instead of jumping *over* the body which would require testing for the opposite, we jump to the true case. We do this because we don't want to have multiple print statements and it lets us fall through the following conditions. We would need multiple print statements because failure for the first expression *is not* failure for the entire expression. Here's how it would look otherwise:

```

1     if (a <= 10)
2         goto check_and;
3     printf("true\n");
4     goto done_if;
5 check_and:
6     if (c <= 100)
7         goto done_if;
8     if (b < c)
9         goto done_if;
10
11     printf("true\n");
12
13 done_if:
14     b |= 0xAA;

```

That is harder to read and has both an extra print and an extra jump.

So let's convert the better version to MIPS (a,b,c = \$t0, \$t1, \$t2):

```

1 .data
2 true_str: .asciiz "true\n"
3
4 .text
5     li    $t8, 10    # get the necessary literals in some unused regs
6     li    $t9, 100
7
8     bgt   $t0, $t8, do_true    # if (a > 10) goto do_true
9     ble   $t2, $t9, done_if    # if (c <= 100) goto done_if
10    blt   $t1, $t2, done_if    # if (b < c) goto done_if
11
12 do_true:
13     li    $v0, 4            # print string
14     la    $a0, true_str     # address of str in a0

```

```

15     syscall
16
17 done_if:
18     ori    $t1, $t1, 0xAA    # b |= 0xAA

```

If-Else Chain

Ok, let's look at a larger example. Say you're trying to determine a student's letter grade based on their score. We're going to need a chain of `if-else-if`'s to handle all cases. Assume `score` is declared and set somewhere before.

```

1  #include <stdio.h>
2
3  int main()
4  {
5      int score;
6      char letter_grade;
7      printf("Enter your score: ");
8      scanf("%d", &score);
9      if (score >= 90) {
10         letter_grade = 'A';
11     } else if (score >= 80) {
12         letter_grade = 'B';
13     } else if (score >= 70) {
14         letter_grade = 'C';
15     } else if (score >= 60) {
16         letter_grade = 'D';
17     } else {
18         letter_grade = 'F';
19     }
20     printf("You got a %c\n", letter_grade);
21     return 0;
22 }

```

With chains like these, if you follow everything we've learned, it comes out looking like this (assuming `score` is `$t0` and `letter_grade` is `$t1`):

```

1  .data
2  prompt:    .asciiz "Enter your score: "
3  grade_str: .asciiz "You got a "
4
5  .text
6  main:
7      li     $v0, 4    # print str
8      la     $a0, prompt
9      syscall
10
11     li     $v0, 5    # read int

```

```

12  syscall
13
14  move    $t0, $v0 # move score into t0
15  li      $t1, 70  # letter_grade default to 'F' ascii value
16
17  li      $t2, 90
18  blt     $t0, $t2, not_a # if (score < 90) goto not_a
19  li      $t1, 65      # leter_grade = 'A'
20  j       grade_done
21
22 not_a:
23  li      $t2, 80
24  blt     $t0, $t2, not_b # if (score < 80) goto not_b
25  li      $t1, 66      # leter_grade = 'B'
26  j       grade_done
27
28 not_b:
29  li      $t2, 70
30  blt     $t0, $t2, not_c # if (score < 70) goto not_c
31  li      $t1, 67      # leter_grade = 'C'
32  j       grade_done
33
34 not_c:
35  li      $t2, 60
36  blt     $t0, $t2, grade_done # if (score < 60) goto grade_done
37  li      $t1, 68      # leter_grade = 'D'
38
39 grade_done:
40  li      $v0, 4       # print str
41  la      $a0, grade_str
42  syscall
43
44  li      $v0, 11      # print character
45  move    $a0, $t1     # char to print
46  syscall
47
48  li      $a0, 10      # print '\n'
49  syscall
50
51  li      $v0, 10      # exit
52  syscall

```

You can see how we set a default value and then test for the opposite of each condition to jump to the next test, until we get one that fails (aka was true in the original C condition) and set the appropriate grade.

You can arrange chains like this in either direction, it doesn't have to match the order of the C code. As long as it works the same, do whatever makes the code simpler and more sensible to you.

Conclusion

Branching and logic and learning to translate from higher level code to assembly is something that takes a lot of practice, but eventually it'll become second nature. We'll get more practice in the chapter on looping which naturally also involves branching.

One final note, there's really no reason to use the `slt` family of opcodes *unless* your professor requires it, ie he says you can't use pseudoinstructions so you're left with `beq`, `bne`, `j` and the `slt` ops. I'll show how you can code without using pseudoinstructions in a later chapter.

Exercises

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

1. Convert the following C code to MIPS.

```
1 #include <stdio.h>
2
3 int main()
4 {
5     int num;
6     printf("Enter an integer: ");
7     scanf("%d", &num);
8
9     if (num > 50) {
10         puts("The number is greater than 50");
11     } else if (num < 50) {
12         puts("The number is less than 50");
13     } else {
14         puts("You entered 50!");
15     }
16
17     return 0;
18 }
```

2. Prompt for the user's name, then tell them whether their name starts with a letter from the first or second half of the alphabet. Be sure to handle both upper and lower case correctly, but assume they entered a valid letter.

[1] https://inst.eecs.berkeley.edu/~cs61c/resources/MIPS_Green_Sheet.pdf

[2] https://en.wikipedia.org/wiki/De_Morgan%27s_laws

Chapter 4: Loops

"Insanity is doing the same thing over and over again and expecting different results."

— Unknown, Often misattributed to Albert Einstein

Before we get into the MIPS, I want to cover something that may be obvious to some but may have never occurred to others. Any loop structure can be converted to any other (possibly with the addition of an `if` statement). So a `for` can be written as a `while` and vice versa. Even a `do-while` can be written as a `for` or `while` loop. Let's look at some equivalencies.

```
1  for (int i=0; i<a; i++) {
2      do_something;
3  }
4
5  int i = 0;
6  while (i < a) {
7      do_something;
8      i++;
9  }
10
11 int i = 0;
12 if (i < a) {
13     do {
14         do_something;
15         i++;
16     } while (i < a);
17 }
18 // you could also have an if (i >= a) goto loop_done; to jump over do-while
```

In general, when writing assembly, it can help to think more in terms of `while` or `do-while` rather than `for` because the former more closely resemble what the assembly looks like in terms of what goes where. Like in the last chapter, where we would think of the `if-else` statements in "jump-form" or "branch-form", we can do the same here, converting `for` to `while` in our head as an intermediary step before going to assembly.

Speaking of "jump-form", let's apply it to the loop above:

```
1  int i=0;
2  if (i >= a)
3      goto done_loop;
4 loop:
5  do_something;
6  i++;
7  if (i < a)
8      goto loop;
```

```
9
10 done_loop:
```

You can see how that starts to look more like assembly. Another thing to note is that unlike with `if` statements where we test for the opposite to jump over the block of code, when you're doing the loop test at the bottom like with a `do-while`, it is unchanged from C because you are jumping to *continue* the loop. If you put the test at the top it becomes inverted, and you put an unconditional jump at the bottom:

```
1   int i=0;
2 loop:
3   if (i >= a)
4       goto done_loop;
5   do_something;
6   i++
7   goto loop:
8
9 done_loop:
```

In general it's better to test at the bottom, both because the condition matches the higher level form, and because when you know the loop is going to execute at least once it requires only one jump + label, rather than 2 since you can forgo the initial `if` check:

```
1   for (int i=0; i<10; i++)
2       do_something;
3
4   // becomes
5
6   int i=0;
7 loop:
8   do_something;
9   i++
10  if (i < 10)
11      goto loop;
```

Ok, now that we've got the theory and structure out of the way, let's try doing a simple one in MIPS.

```
1   int sum = 0;
2   for (int i=0; i<100; i++) {
3       sum += i;
4   }
```

That's about as basic as it gets, adding up the numbers 0 to 99.

```
1   li    $t0, 0    # sum = 0
2   li    $t1, 1    # i = 1 we can start at 1 because obviously adding 0 is
```

```

pointless
3    li    $t2, 100
4 loop:
5    addi  $t0, $t0, $t1    # sum += i
6    addi  $t1, $t1, 1      # i++
7    blt   $t1, $t2, loop   # while (i < 100)

```

Ok I don't think there's much point in doing any more without getting to what loops are most often used for, looping through data structures, most commonly arrays.

Looping Through Arrays

Looping and arrays go together like peanut butter and jam. An array is a sequence of variables of the same type, almost always related in some way. Naturally, you want to operate on them all together in various ways; sorting, searching, accumulating, etc. Given that the only way to do that is with loops, in this section we'll cover different ways of looping through arrays, including multidimensional arrays.

1D Arrays

Let's pretend there's an array `int numbers[10];` filled with 10 random numbers.

```

1    int total = 0;
2    for (int i=0; i<10; i++) {
3        total += numbers[i];
4    }

```

There are several ways to do this. The first is the most literal translation.

```

1    li    $t0, 0    # total = 0
2    li    $t1, 0    # i = 0
3    la    $t2, numbers    # t2 = numbers
4    li    $t3, 10
5 sum_loop:
6    sll   $t4, $t1, 2    # t4 = i*sizeof(int) == i*4
7    add   $t4, $t4, $t2   # t4 = &numbers[i]
8    lw    $t4, 0($t4)     # t4 = numbers[i]
9    add   $t0, $t0, $t4    # total += numbers[i]
10
11    addi  $t1, $t1, 1    # i++
12    blt   $t1, $t3, sum_loop    # while (i < 10)

```

We initialize the relevant variables beforehand (`numbers` and `10` could be loaded every iteration but that's less efficient). Now what's with the `i*4`? We already discussed using shifts to multiply and divide by powers of 2 in a previous chapter, but here we're doing something that higher level languages do automatically for you every time you do an array access. When you access the *i*'th element, under the hood it is multiplying *i* by the size of the type of the array and adding that

number of bytes to the base address and then loading the element located there.

If you're unfamiliar with the C syntax in the comments, `&` means "address of", so `$t4` is being set to the address of the `i`'th element. Actually that C syntax is redundant because the `&` counteracts the brackets. In C adding a number to a pointer does pointer math (ie it multiplies by the size of the items as discussed above). This means that the following comparison is true:

```
&numbers[i] == numbers + i
```

which means that this is true too

```
&numbers[0] == numbers
```

The reason I use the form on the left in C/C++ even when I can use the right is it makes it more explicit and obvious that I'm getting the address of an element of an array. If you were scanning the code quickly and saw the expression on the right, you might not realize that's an address at all, it could be some mathematical expression (though the array name would hopefully clue you in if it was picked well).

Anyway, back to the MIPS code. After we get the address of the element we want, we have to actually read it from memory (ie load it). Since it's an array of words (aka 4 byte ints) we can use load word, `lw`.

Finally, we add that value to `total`, increment `i`, and perform the loop check.

Now, I said at the beginning that this was the most literal, direct translation (not counting the restructuring to a `do-while` form). However, it is not my preferred form because it's not the simplest, nor the shortest.

Rather than calculate the element address every iteration, why not keep a pointer to the current element and iterate through the array with it? In C what I'm suggesting is this:

```
1  int* p = &numbers[0];
2  int i = 0, total = 0;
3  do {
4      total += *p;
5      i++;
6      p++;
7  } while (i < 10);
```

In other words, we set `p` to point at the first element and then increment it every step to keep it pointing at `numbers[i]`. Again, all mathematical operations on pointers in C deal in increments of the byte size of the type, so `p++` is really adding `1*sizeof(int)`.

```
1  li    $t0, 0      # total = 0
2  li    $t1, 0      # i = 0
3  la    $t2, numbers # p = numbers
4  li    $t3, 10
5  sum_loop:
6  lw    $t4, 0($t2)  # t4 = *p
```

```

7    add    $t0, $t0, $t4 # total += *p
8
9    addi   $t1, $t1, 1    # i++
10   addi   $t2, $t2, 4    # p++ ie p += sizeof(int)
11   blt    $t1, $t3, sum_loop # while (i < 10)

```

Now, that may not look much better, we only saved 1 instruction, and if we were looping through a string (aka an array of characters, `sizeof(char) == 1`) we wouldn't have saved any. However, imagine if we weren't using `sll` to do the multiply but `mult`. That would take 3 instructions, not 1. Even `mul` would take 2. Remember we *would* have to use one of those if we were iterating through an array of structures with a size that wasn't a power of 2.

However, there is one more variant that you can use that can save a few more instructions. Instead of using `i` and `i < 10` to control the loop, use `p` and the address just past the end of the array. In C it would be this:

```

1    int* p = &numbers[0];
2    int* end = &numbers[10];
3    int total = 0;
4    do {
5        total += *p;
6        p++;
7    } while (p < end);

```

You could also use `!=` instead of `<`. This is similar to using the `.end()` method on many C++ data structures when using iterators. Now the MIPS version:

```

1    li     $t0, 0        # total = 0
2    la     $t2, numbers   # p = numbers
3    addi   $t3, $t2, 40    # end = &numbers[10] = numbers + 10*sizeof(int)
4 sum_loop:
5    lw     $t4, 0($t2)     # t4 = *p
6    add    $t0, $t0, $t4   # total += *p
7
8    addi   $t2, $t2, 4     # p++ ie p += sizeof(int)
9    blt    $t2, $t3, sum_loop # while (p < end)

```

So we dropped from 10 to 7 instructions and even more if we had had to do `mul` or `mult` originally. And this was for a 1D array. Imagine if you had 2 or 3 indices you had to use to calculate the correct offset. That's what we go over in the next section.

2D Arrays

The first thing to understand is what's really happening when you declare a 2D array in C. The contents of a 2D array are tightly packed, in row-major order, meaning that all the elements from the first row are followed by all the elements of the second row and so on. What this means is that a 2D array is equivalent to a 1D array with `rows*cols` elements in the same order:

```

1  #define ROWS 2
2  #define COLS 4
3  // The memory of these two arrays are identical
4  int array[ROWS][COLS] = { { 1, 2, 3, 4 }, { 5, 6, 7, 8 } };
5  int array1d[ROWS*COLS] = { 1, 2, 3, 4, 5, 6, 7, 8 };

```

See the code example [2d_arrays.c](#) for more details.

What this means is that when we declare a 2D array, it's basically a 1D array with the size equal to rows * columns. Also, when we loop through a 2D array, we can often treat it like a 1D array with a single loop. So everything that we learned before applies.

Let's do an example.

```

1  for (int i=0; i<rows; i++) {
2      for (int j=0; j<cols; ++j) {
3          array[i][j] = i + j;
4      }
5  }
6
7  // becomes
8
9  int r, c;
10 for (int i=0; i<rows*cols; i++) {
11     r = i / cols;
12     c = i % cols;
13     array[i] = r + c;
14 }

```

So assuming `rows` and `cols` are in `$a0` and `$a1` (and nonzero), it would look like this:

```

1  la    $t0, array    # p = &array[0]
2  li    $t1, 0        # i = 0
3  mult  $a0, $a1      # a0 = rows, a1 = cols
4  mflo  $t2           # t2 = rows * cols
5 loop:
6  div   $t1, $a1      # r = i / cols
7  mflo  $t3           # c = i % cols
8  mfhi  $t4           # c = i % cols
9  add   $t3, $t3, $t4  # t3 = r + c
10
11 sw    $t3, 0($t0)    # array[i] = *p = r + c
12
13 addi  $t1, $t1, 1    # i++
14 addi  $t0, $t0, 4     # p++ (keep pointer in sync with i, aka p = &array[i])
15 blt   $t1, $t2, loop # while (i < rows*cols)

```

You might ask if it's worth it to convert it to a single loop when you still need the original *i* and *j* as if you were doing nested loops. Generally, it is much nicer to avoid nested loops in assembly if you can. There are many cases when you get the best of both worlds though. If you're doing a clear for example, setting the entire array to a single value, there's no need to calculate the row and column like we did here. I only picked this example to show how you could get them back if you needed them.

For comparison here's the nested translation (while still taking advantage of the 1D arrangement of memory and pointer iterators):

```

1  la    $t0, array    # p = &array[0]
2  li    $t1, 0        # i = 0
3 looprows:
4  li    $t2, 0        # j = 0
5 loopcols:
6  add   $t3, $t1, $t2  # t3 = i + j
7  sw    $t3, 0($t0)    # array[i][j] = *p = i + j
8
9  addi  $t2, $t2, 1    # j++
10 addi  $t0, $t0, 4     # p++ (keep pointer in sync with i and j, aka p =
    &array[i][j])
11 blt   $t2, $a1, loopcols # while (j < cols)
12
13 addi  $t1, $t1, 1    # i++
14 blt   $t1, $a0, looprows # while (i < rows)

```

It's a bit shorter but again, how much are the extra label and branch worth? For me, this one's a toss up. On the other hand, either of the last 2 versions are better than the literal translation below:

```

1  la    $t0, array    # p = &array[0]
2  li    $t1, 0        # i = 0
3 looprows:
4  li    $t2, 0        # j = 0
5 loopcols:
6  add   $t3, $t1, $t2  # t3 = i + j
7
8  # need to calculate the byte offset of element array[i][j]
9  mult  $t1, $a1
10 mflo  $t4            # i * cols
11
12 add   $t4, $t4, $t2  # t4 = i * cols + j
13
14 sll   $t4, $t4, 2    # t4 = (i * cols + j) * sizeof(int)
15 add   $t4, $t4, $t0  # t4 = &array[i][j] (calculated as array + (i*cols +
    j)*4)
16
17 sw    $t3, 0($t4)    # array[i][j] = i + j
18
19 addi  $t2, $t2, 1    # j++

```

```

20    blt    $t2, $a1, loopcols    # while (j < cols)
21
22    addi   $t1, $t1, 1           # i++
23    blt    $t1, $a0, looprows    # while (i < rows)

```

That chunk in the middle calculating the offset of every element? Not only is it far slower than iterating the pointer through the array, but you can imagine how much worse it would be for a 3D array with 3 nested loops.

Conclusion

Hopefully after those examples you have a more solid understanding of looping in MIPS and how to transform various loops and array accesses into the form that makes your life the easiest. There is more we could cover here, like looping through a linked list, but I think that's beyond the scope of what we've covered so far. Perhaps in a later chapter.

Exercises

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

1. Convert the following C code to MIPS. If using SPIM, you can just hard code a "random" number between 0 and 100.

```

1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main()
5  {
6      int num = rand() % 101
7      int guess;
8      puts("I'm thinking of a number 0-100. Try to guess it!");
9      while (1) {
10         printf("Guess a number: ");
11         scanf("%d", &guess);
12
13         if (guess > num) {
14             puts("Too high!");
15         } else if (guess < num) {
16             puts("Too low!");
17         } else {
18             break;
19         }
20     }
21
22     printf("Correct, it was %d!\n", num);
23
24     return 0;

```

```
25 }
```

2. Write a MIPS program to find and print the average of the following array. Use integer division.

```
1 .data
2 array: .word 93,8,78,-6,51,49,3,2,128,0
```

3. Write a program to find the min and max of the array in the previous exercise

Chapter 5: Functions and the MIPS Calling Convention

While I'm sure everyone here probably knows what functions are, you might be wondering what a "Calling Convention" is. In short, it is an agreement between the caller and callee about how to treat/use certain registers. We'll get to the why and how later.

Functions

In assembly, a function is simply a label with a return instruction associated with it; because this is far more ambiguous than a function in a higher level language, it is good practice to only have a single return instruction associated with a function.^[1] A comment above the label is also helpful. Together those help you quickly see the start and end of the function.

```
1 void func1() {}
```

would be

```
1 # void func1()
2 func1:
3     # body goes here
4     jr     $ra
```

As you can see my policy is to put a single line comment of the C prototype above label.

But how do you call a function in assembly? You use the instruction Jump and Link: `jal func_label`. Let's change the hello world program from chapter 0 to call a function:

```
1 .data
2 hello: .asciiz "Hello World!\n"
3
4 .text
5 main:
6     jal hello_world
7
8     li $v0, 10    # exit syscall
9     syscall
10
11
12 # void hello_world()
13 hello_world:
14     li $v0, 4      # print string system call
15     la $a0, hello  # load address of string to print into a0
16     syscall
17
```

What `jal` actually does, is save the address of the next instruction to `$ra` and then do an unconditional jump to the function label. So you could achieve the same results with the following:

```
jal    func

# is equivalent to

la     $ra, next_instr
j      func
next_instr:
```

That would get tiring and ugly fast though, having to come up with unique labels for the next instruction every time. You also might be confused about why the greensheet says `jal` saves PC+8 in `$ra` instead of PC+4. The reason is that MIPS technically has delayed branching, i.e. a single instruction after every branch instruction is executed before the jump actually happens. So `jal` adds 8 instead of 4 to account for that extra instruction delay. However, every simulator we've mentioned does non-delayed branching by default so you can ignore it.

The Convention

We've gone as far as we can without starting to talk about registers and their purposes in functions. You can think of registers as variables^[2] that are part of the CPU. In this case, since we're dealing with a 32-bit MIPS architecture, they are 32-bit (aka 4 bytes, 1 word) variables. Since they're part of the CPU, they exist for the life of the program and the whole program shares the same registers.

But how does that work? If all parts of the program use the same 32 registers, how does one function not stomp all over what another was doing when it uses them? In fact, how do functions communicate at all? How do they pass arguments or return results? All these questions are solved by deciding on a "Calling Convention". It's different for different architectures and even different operating systems on the same architecture. This is because different architectures have different numbers of registers, and some registers like `$ra` have hardcoded uses. The op `jal` modifies `$ra`, and `$0` is a constant 0 and there's no way to change either of those facts. That still leaves a lot of flexibility when designing a calling convention. While they mostly match, you'll find several variations of MIPS calling conventions online. They usually differ in how they setup a stack frame. The convention covered in this chapter is consistent with, and sufficient for, almost every college course I've ever heard of.

Regardless, what matters is that the calling convention works by setting rules (and guidelines) for register use, and when/how to use the stack.

If you're unfamiliar with the runtime stack, it's exactly what it sounds like. It's a Last-In-First-Out (LIFO) data structure that you can use to store smaller values in a program. It grows in a negative direction, so to allocate 12 bytes, you would subtract 12 from the stack pointer (in MIPS that's `$sp`).

MIPS specifically designates certain registers to be used for passing arguments (at least the first 4),

others for return values, and others for misc. temporary or saved values. The rest are special use registers like `$ra`.

The quickest way to summarize is to look at the table on the greensheet which is reproduced (with some modifications) below:

Table 4. MIPS Registers and Uses

Name	Number	Use	Preserved Across a Call
<code>\$zero</code>	0	Constant 0	N.A.
<code>\$at</code>	1	Assembler Temporary (used to expand pseudo-ops)	No
<code>\$v0-\$v1</code>	2-3	Function Results and Expression Evaluation	No
<code>\$a0-\$a3</code>	4-7	Arguments	No
<code>\$t0-\$t7</code>	8-15	Temporaries	No
<code>\$s0-\$s7</code>	16-23	Saved Temporaries	Yes
<code>\$t8-\$t9</code>	24-25	Temporaries	No
<code>\$k0-\$k1</code>	26-27	Reserved for OS Kernel	No
<code>\$gp</code>	28	Global Pointer	Yes
<code>\$sp</code>	29	Stack Pointer	Yes
<code>\$fp</code> (or <code>\$s8</code>)	30	Frame Pointer if necessary or can be another saved reg	Yes
<code>\$ra</code>	31	Return Address	No

To summarize, you have 16 registers that can be used anytime for temporary values, though some have special uses too (the v, a, and t registers). You have 8 s registers that have to be saved on the stack if you use them, plus `$ra` as well. The `$zero` register is obviously a special case.

The `$sp` register is technically preserved but not in the same way. Basically what you allocate (subtract) you have to deallocate (add) before returning from a function, thus preserving the original value.

You can ignore `$at`, `$k0-$k1`, `$gp` and most of the time `$fp` too. In over 7 years of tutoring I've helped students with MIPS from at least 2 dozen different colleges and I think I've only seen a professor force his students to use `$fp` or pass more than 4 arguments twice. You can see^[3] register 30 sometimes referred to as `$s8` rather than, or in addition to, `$fp` which shows you how rarely it's actually used/needed as a frame pointer.

Basic example

Let's start with something simple that doesn't use the stack.

```
int hello_name_number(char* name, int number)
```

```

{
    printf("Hello %s!\n", name);
    return number + 10;
}

```

According to the convention that becomes:

```

1 .data
2 hello_space: .asciiz "Hello "
3 exclaim_nl: .asciiz "!\n"
4
5 .text
6 # int hello_name_number(char* name, int number)
7 hello_name_number:
8     move    $t0, $a0    # save name in t0 since we need a0 for the syscall
9
10    li      $v0, 4        # print string
11    la      $a0, hello_space
12    syscall
13
14    move    $a0, $t0    # print name (v0 is still 4)
15    syscall
16
17    la      $a0, exclaim_nl # print "!\n"
18    syscall
19
20
21    addi    $v0, $a1, 10 # return number + 10
22    jr      $ra

```

Some things to note, syscalls are not function calls so we can "save" `$a0` in a `t` register and know that it'll still be there when the syscall is done. In the same way, we know that `$v0` is still the same so we don't have to keep setting it to 4 for print string. Lastly, to return a value, we make sure that value is in `$v0` before returning.

Using the Stack

First, let's establish the rules on when you *have* to use the stack (You can always use it for arbitrary local variables, like a local array for example, but generally don't if you don't have a good reason).

1. You call another function, ie you're a non-leaf function.

This means you have to save `$ra` on the stack at the very least, otherwise when you do your `jr $ra` you'd jump back into yourself (right after the last `jal` instruction). This does not apply to main because you don't/shouldn't return from main, you should call the exit (or exit2) syscall (10 or 17).

2. You need to save values across a function call (automatically includes reason 1).

This is fairly common for non-trivial functions. Obvious examples are calling a function in a loop or loops (you'd have to preserve the iterator(s)), and many recursive functions.

3. You run out of temporary registers and overflow into the s registers.

This is very rare. The most common reason this "happens" is people forget they have 10 t registers instead of 8 like s registers and even if they remember that they forget they can also use the a and v registers for temporaries. 16 is more than enough to handle pretty much any function because you rarely need 17 discrete values at the same time.

Let's look at an example for the first two. Any example for the last rule would be prohibitively large and complicated.

```
1 int non_leaf()  
2 {  
3     func1();  
4     return 42  
5 }
```

This calls the empty function discussed at the top of this chapter.

```
1 #int non_leaf()  
2 non_leaf:  
3     addi    $sp, $sp, -4 # space to save 1 register, $ra  
4     sw      $ra, 0($sp) # store $ra in the newly allocated stack space  
5  
6     jal     func1  
7  
8     li      $v0, 42      # return 42  
9  
10    lw      $ra, 0($sp) # restore original $ra  
11    addi    $sp, $sp, 4  # pop the stack  
12    jr      $ra
```

The bit of code at the top and bottom of the function are called the prologue and epilogue respectively for obvious reasons. We allocate 4 bytes on the stack by subtracting 4 (I add a negative rather than subtract because I can copy-paste the line with a single character change for the epilogue). Then we store the current `$ra` in that space at the new top of the stack. Then before we exit we have to load it back and pop the stack.

If we didn't save and restore `$ra` we would jump to line 7 when we do our `jr $ra` and then we'd be in an infinite loop.

Next we have the second case, where we need to preserve regular local values across a function call.

```
1 void print_letters(char letter, int count)  
2 {
```

```

3   for (int i=0; i<count; i++) {
4       putchar(letter);
5   }
6   putchar('\n');
7 }
8
9 int save_vals()
10 {
11     for (int i=0; i<10; i++) {
12         print_letters('A'+i, i+1);
13     }
14     return 8;
15 }

```

That becomes this in mips:

```

1 #void print_letters(char letter, int count)
2 print_letters:
3     ble    $a1, $0, exit_pl    # if (count <= 0) goto exit_pl
4     li     $v0, 11             # print character
5 pl_loop:
6     syscall
7     addi   $a1, $a1, -1        # count--
8     bgt    $a1, $0, pl_loop    # while (count > 0)
9
10    li     $a0, 10             # '\n'
11    syscall
12
13 exit_pl:
14     jr     $ra
15
16
17 #int save_vals()
18 save_vals:
19     addi   $sp, $sp, -12
20     sw     $ra, 0($sp)
21     sw     $s0, 4($sp)
22     sw     $s1, 8($sp)
23
24     li     $s0, 0    # i = 0
25     li     $s1, 10
26 sv_loop:
27     addi   $a0, $s0, 65    # i + 'A'
28     addi   $a1, $s0, 1     # i + 1
29     jal    print_letters
30
31     addi   $s0, $s0, 1     # i++
32     blt    $s0, $s1, sv_loop    # while (i < 10)
33
34     lw     $ra, 0($sp)

```

```

35     lw     $s0, 4($sp)
36     lw     $s1, 8($sp)
37     addi   $sp, $sp, 12
38     jr     $ra

```

Notice that for `print_letters`, we not only convert the loop to a `do-while`, but we also use the parameter `count` as the iterator to count *down* to 0. It saves us an instruction initializing an `i`.

Second, for `save_vals`, we save not only `$ra` because we call another function, but also two `s` registers to save `i` and our stopping point. The second is not actually necessary; because it's a constant, we could load 10 into a register right before the check every iteration of the loop. Which version is better depends on several factors, like how long or complex the loop is, how many times it executes, and of course personal preference.

Recursive Functions

Let's do a classic recursive function, the fibonacci sequence.

```

1 int fib(int n)
2 {
3     if (n <= 1)
4         return n;
5
6     return fib(n-2) + fib(n-1);
7 }

```

You can see how, at the very least, we'll have to save `$ra` and `n`, because we need the original even after the first recursive call. It's not as obvious, but we'll also have to save the return value of the first call so we'll still have it to do the addition after the second. You might think this would require using two `s` regs, but does it? Let's see...

```

1 #int fib(int n)
2 fib:
3     addi   $sp, $sp, -8
4     sw     $ra, 0($sp)
5     sw     $s0, 4($sp)
6
7     move   $v0, $a0      # prepare to return n
8     li     $t0, 1
9     ble    $a0, $t0, exit_fib # if (n <= 1) goto exit_fib (ie return n)
10
11     move   $s0, $a0      # save n
12
13     addi   $a0, $a0, -2   # a0 = n - 2
14     jal    fib           # fib(n-2)
15
16     addi   $a0, $s0, -1   # a0 = n - 1, prep arg first so we can use s0 to save
v0

```

```

17  move    $s0, $v0      # save return of fib(n-2) in s0
18  jal     fib           # fib(n-1)
19
20  add     $v0, $v0, $s0  # v0 = fib(n-1) + fib(n-2)
21
22  exit_fib:
23  lw      $ra, 0($sp)
24  lw      $s0, 4($sp)
25  addi    $sp, $sp, 8
26  jr      $ra

```

Notice how we don't have to save `n` any sooner than necessary, ie right before we have to use `$a0` to setup the first recursive call. Also, the ordering of lines 16 and 17 is important. We needed the original `n` to calculate `n-1` but once that's in `$a0` ready for the call, because we won't need `n` again afterward, we can now use `$s0` to preserve the return value of the first call.

Some of you, if you were paying attention, might point out that you could save a few instructions of performance if you moved the base case testing before the prologue as long as you put the exit label after the epilogue. This is true, but I'd recommend against it unless you were really trying to eke out every last microsecond. It's nicer/cleaner to keep the prologue and epilogue as the first and last things; they're one more thing to catch your eye and help delineate where functions start and end. Regardless, if you're curious, you can see that version, along with every other function in this chapter in the included program [calling.s](#).

Conclusion

While grasping the basics of a calling convention is not too difficult, it takes practice to get used to it. There are many things that we haven't covered in this chapter, like how to pass more than 4 arguments, or use `$fp`, or handle floating point arguments or return values. The latter at least, will be covered in the next chapter.

Exercises

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

1. Implement the following functions in MIPS and write a program to demonstrate their use. You can reuse much of your code from the previous chapter's exercises.

```

1 // return the min or max
2 int get_min(int* array, int size);
3 int get_max(int* array, int size);
4
5 // return the index of the min/max
6 int locate_min(int* array, int size);
7 int locate_max(int* array, int size);
8
9 // return the average of the array

```

```
10 int calc_average(int* array, int size);
```

2. The Collatz conjecture is defined as follows: start with any positive integer n, if n is even, divide by 2, otherwise, multiply by 3 and add 1. The conjecture is that all sequences will eventually reach 1. Write 2 versions of the collatz function, one iterative and one recursive. Print out the sequence as they go.

```
1 void collatz_iterative(int n);  
2 void collatz_recursive(int n);  
3  
4 // for an added challenge return the number of steps taken to reach 1  
5 // you can remove the printing of the steps  
6 int collatz_iterative2(int n);  
7 int collatz_recursive2(int n);
```

[1] I do not agree with an ironclad "one return" policy in higher level languages. Sometimes returning early results in cleaner code, sometimes not. Similarly, `goto` is not evil and there are rare cases where using it creates the best code.

[2] Obviously the zero register is not really a variable. I never understood how people could say "const variable" with a straight face, it's literally an oxymoron.

[3] It's an [old link](#), but not as old as SPIM so maybe using it for a frame pointer was added later

Chapter 6: Floating Point Types

Up to this point we haven't really mentioned floating point values or instructions at all, except how to declare them in the `.data` section and the syscalls for reading and printing them. There are two reasons we've left them alone till now. First, they use a whole separate set of registers and instructions. Second, and partly because of the first reason, most MIPS college courses do not ever require you to know or use floating point values. Since this book is targeted at college students, if you know you won't need to know this feel free to skip this chapter.

Floating Point Registers and Instructions

While the greensheet contains a nice table for the normal registers it is completely lacking for the floating point registers. There are 32 32-bit floating point registers. You can use them all for floats but they are paired even-odd for doubles. In other words, you can only use even numbers for doubles, because storing a double at `$f0` actually uses `$f0` and `$f1` because it takes 64 bits/8 bytes.

As far as the calling conventions for floating point registers, it is actually [hard](#) to [find](#) anything definitive and clear even for the basics. You could make up your own but the float/double syscalls, and the tiny code snippet in [Patterson and Hennessy](#) were at least consistent with this old website^[1] so we'll go with that. I have seen at least one course page where the prof wanted *all* float registers preserved which seems excessive and ridiculous but prof's are gonna prof.

Table 5. MIPS Floating Point Registers and Uses

Name	Use	Preserved Across a Call
<code>\$f0-\$f2</code>	Function Results	No
<code>\$f4-\$f10</code>	Temporaries	No
<code>\$f12-f14</code>	Arguments	No
<code>\$f16-f18</code>	Temporaries	No
<code>\$f20-f30</code>	Saved Temporaries	Yes

This table is based on doubles so it may look like it's skipping odd registers but they're included where the even they're paired with is. So, for example you actually have 4 registers for float arguments `$f12` through `$f15` but only 2 for doubles `$f12` and `$f14`. Similarly you have 12 saved registers for floats but 6 for doubles.

You might find things like [this](#), which seems to say that SPIM doesn't support using the odd registers at all but both example programs for this chapter use `$f1` and work with both SPIM and MARS. Given that, and the fact that it references instructions `[l.d]` and `[l.s]` which don't work (`li.s` and `li.d` do, see below), it's probably really out of date.

Most of the next table is actually on the Greensheet but not all of it and I thought it worth reproducing here.

Table 6. MIPS floating point instructions (and pseudoinstructions)

Name	Opcode	Format	Operation
Load Word to Coprocessor 1	lwc1 (or l.s)	lwc1 ft, n(rs)	$F[ft] = M[R[rs]+n]$
Store Word from Coprocessor 1	swc1 (or s.s)	swc1 ft, n(rs)	$M[R[rs]+n] = F[ft]$
Load Double to Coprocessor 1	ldc1 (or l.d)	ldc1 ft, n(rs)	$F[ft] = M[R[rs]+n]$ $F[ft+1] = M[R[rs]+n+4]$
Store Double from Coprocessor 1	sdc1 (or s.d)	sdc1 ft, n(rs)	$M[R[rs]+n] = F[ft]$ $M[R[rs]+n+4] = F[ft+1]$
Move From Coprocessor 1	mfc1	mfc1 rd, fs	$R[rd] = F[fs]$
Move To Coprocessor 1	mtc1	mtc1 rd, fs	$F[fs] = R[rd]$
Convert Word To Single Precision	cvt.s.w	cvt.s.w fd, fs	$F[fd] = (\text{float})F[fs]$
Convert Single Precision To Word	cvt.w.s	cvt.w.s fd, fs	$F[fd] = (\text{int})F[fs]$
Convert Word To Double Precision	cvt.d.w	cvt.d.w fd, fs	$F[fd] = (\text{double})F[fs]$
Convert Double Precision To Word	cvt.w.d	cvt.w.d fd, fs	$F[fd] = (\text{int})F[fs]$
Branch on FP True	bc1t	bc1t label	if (FPcond) goto label;
Branch on FP False	bc1f	bc1f label	if (!FPcond) goto label;
FP Compare	c.y.x	c.y.x fs, ft	$FPcond = (F[fs] \text{ op } F[ft]) ? 1 : 0$
Absolute Value	abs.x	abs.x fs, ft	$F[fs] = (F[ft] > 0) ? F[ft] : -F[ft]$
Add	add.x	add.x fd, fs, ft	$F[fd] = F[fs] + F[ft]$
Subtract	sub.x	sub.x fd, fs, ft	$F[fd] = F[fs] - F[ft]$
Multiply	mul.x	mul.x fd, fs, ft	$F[fd] = F[fs] * F[ft]$
Divide	div.x	div.x fd, fs, ft	$F[fd] = F[fs] / F[ft]$
Negation	neg.x	neg.x fs, ft	$F[fs] = -F[ft]$
Move	mov.x	mov.x fd, fs	$F[fd] = F[fs]$

With all of the opcodes that end in .x, the x is either s for single precision or d for double precision.

The y in the Compare instructions are one of eq, lt, le. Naturally *op* would be the matching ==, <, <=. Unfortunately, you don't get *not equal*, *greater than*, or *greater equal*, even as pseudoinstructions, but it's easy enough to flip the order of operands or branch on the opposite result.

Practice

We're going to briefly go over some of the more different aspects of dealing with floating point numbers, but since most of it is the same but with a new set of registers and calling convention, we

won't be rehashing most concepts.

Getting Floating Point Literals

The first thing to know when dealing with floats is how to get float (or double) literals into registers where you can actually operate on them.

There are two ways. The first, and simpler way, is to declare them as globals and then use the `lwc1` or `ldw1` instructions:

```
1 .data
2 a:    .float 3.14159
3 b:    .double 1.61
4
5 .text
6 main:
7
8     la    $t0, a
9     lwc1  $f0, 0($t0)    # get a into $f0
10
11     la    $t0, b
12     ldc1  $f2, 0($t0)    # get b into $f2-3
13
14     # other code here
```

The second way is to use the regular registers and convert the values. Of course this means unless you want an integer value, you'd have to actually do it twice and divide, and even that would limit you to rational numbers. It looks like this.

```
1     mtc1  $0, $f0        # move 0 to $f0 (0 integer == 0.0 float)
2
3     # get 4 to 4.0 in $f2
4     li    $t0, 4
5     mtc1  $t0, $f2
6     cvt.s.w $f2, $f2    # convert 4 to 4.0
```

As you can see, other than 0 which is a special case, it requires at least 3 instructions, more than the 2 (or 1 if you load directly from the address) of the first method.

NOTE

There is a 3rd way that is even easier, but it's only supported in SPIM. The pseudoinstructions `li.s` and `li.d` work exactly like `li` except to load float and double literals into float/double registers.

Branching

Branching based on floating point values is slightly different than normal. Instead of being able to

test and jump in a single convenient instruction, you have to test first and then jump in a second instruction if the test was true or not. This is the same way x86 does it. The test sets a special control/flag register (or a certain bit or bits in the register) and then all jumps are based on its state.

Using it looks like this:

```
1  c.lt.s  $f0, $f2  # fpcond = f0 < f2
2  bc1t   was_less  # if (f0 < f2) goto was_less
3
4  # do something for f0 >= f2
5
6  j      blah
7 was_less:
8
9  # do something for f0 < f2
10
11 blah:
```

Functions

Finally, lets do a simple example of writing a function that takes a float and returns a float. I'm not going to bother doing one for doubles because it'd be effectively the same, or doing one that requires the stack, because the only differences from normal are a new set of registers and knowing which ones to save or not from the table above.

So, how about a function to convert a fahrenheit temperature to celsius:

```
1 .data
2
3 # 5/9 = 0.5 with 5 repeating
4 fahrenheit2celsius: .float 0.555555
5
6 .text
7 # float convert_F2C(float degrees_f)
8 convert_F2C:
9     la      $t0, fahrenheit2celsius
10    lwc1     $f0, 0($t0)    # get conversion factor
11
12    # C = (F - 32) * 5/9
13    li      $t0, 32
14    mtc1     $t0, $f1        # move int 32 to f1
15    cvt.s.w  $f1, $f1        # convert to 32.0
16
17
18    sub.s    $f12, $f12, $f1 # f12 = degrees - 32
19
20    mul.s    $f0, $f0, $f12  # f0 = 0.555555 * f12
21
```

You can see we follow the convention with the argument coming in `$f12` and the result being returned in `$f0`. In this function we use both methods for getting a value into float registers; one we load from memory and the other, being an integer, we move and convert.

Conclusion

As I said before, it is rare for courses to even bother covering floating point instructions or assign any homework or projects that use them, but hopefully this brief overview, combined with the knowledge of previous chapters is sufficient.

There are also 2 example programs [conversions.s](#) and [calc_pi.s](#) for you to study.

Exercises

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

1. Write a program to convert an input in minutes to hours.
2. Write the following functions and use them in a program.

```
1 float miles2kilometers(float miles);  
2 float pounds2kilograms(float pounds);
```

[1] Besides, [it's](#), the same one we referenced last chapter about `fp == s8`

Chapter 7: Tips and Tricks

This chapter is a grab bag of things you can do to improve your MIPS programs and make your life easier.

Formatting

You may have noticed I have a general format I like to follow when writing MIPS (or any) assembly. The guidelines I use are the following

1. **1 indent for all code excluding labels/macros/constants.**

I use hard tabs set to a width of 4 but it really doesn't matter as long as it's just 1 indent according to your preferences.

2. **Use *spaces* to align the first operand of all instructions out far enough.**

Given my 4 space tabs, this means column 13+ (due to syscall, though I often stop at 11). The reason to use spaces is to prevent the circumstances that gave hard tabs a bad name. When you use hard tabs for alignment, rather than indentation, and then someone else opens your code with their tab set to a different width, suddenly everything looks terrible. Thus, tabs for indentation, spaces for alignment. Or as is increasingly common (thanks Python), spaces for everything but I refuse to do that to the poor planet.^[1]

3. **A comma and a single space between operands.**

The simulators don't actually require the comma but since other assembly languages/assemblers do, you might as well get used to it. Besides I think it's easier to read with the comma, though that might be me comparing it to passing arguments to a function.

4. **Comment every line or group of closely related lines with the purpose.**

This is often simply the equivalent C code. You can relax this a little as you get more experience.

5. **Use a blank line to separate logically grouped lines of code.**

While you can smash everything together vertically, I definitely wouldn't recommend it, even less than I would in a higher level language.

6. **Put the `.data` section at the top, similar to declaring globals in C.**

There are exceptions for this. When dealing with a larger program with lots of strings, it can be convenient to have multiple `.data` sections with the strings you're using declared close to where you use them. The downside is you have to keep swapping back and forth between `.text` and `.data`.

Misc. General Tips

1. **Try to use registers starting from 0 and working your way up.**

It helps you keep track of where things are (esp. combined with the comments). This can fall apart when you discover you forgot something or need to modify the code later and it's often not worth changing all the registers you're already using so you can maintain that nice sequence. When that happens I'll sometimes just pick the other end of sequence (ie \$t9 or \$s7) since if it's out of order I might as well make it obvious.

2. Minimize your jumps, labels, and especially your level of nested loops.

This was already covered in the chapters on branching and loops but it bears repeating.

3. In your prologue save \$ra first, if necessary, then all s regs used starting at \$s0.

Then copy paste the whole thing to the bottom, move the first line to the bottom and change the number to positive and change all the sw to lw.

```
func:
    addi    $sp, $sp, -20
    sw      $ra, 0($sp)
    sw      $s0, 4($sp)
    sw      $s1, 8($sp)
    sw      $s2, 12($sp)
    sw      $s3, 16($sp)

    # body of func here that calls another function or functions
    # and needs to preserve 4 values across at least one of those calls

    lw      $ra, 0($sp)
    lw      $s0, 4($sp)
    lw      $s1, 8($sp)
    lw      $s2, 12($sp)
    lw      $s3, 16($sp)
    addi    $sp, $sp, 20
```

Constants

One of the easiest things you can do to make your programs more readable is to use defined constants in your programs. Both MARS and SPIM have ways of defining constants similar to how C defines macro constants; ie they aren't "constant variables" that take up space in memory, it's as if a search+replace was done on them right before assembling the program.

Let's look at our Hello World program using constants for SPIM and MARS

SPIM:

```
1 sys_print_str = 4
2 sys_exit = 10
3
4 .data
5 hello: .asciiz "Hello World!\n"
```

```

6
7 .text
8 main:
9     li    $v0, sys_print_str
10    la    $a0, hello # load address of string to print into a0
11    syscall
12
13    li    $v0, sys_exit
14    syscall

```

MARS:

```

1 .eqv sys_print_str 4
2 .eqv sys_exit 10
3
4 .data
5 hello: .asciiz "Hello World!\n"
6
7 .text
8 main:
9     li    $v0, sys_print_str
10    la    $a0, hello # load address of string to print into a0
11    syscall
12
13    li    $v0, sys_exit
14    syscall

```

Macros

MARS supports function style macros that can shorten your code and improve readability in some cases (though I feel it can also make it worse or be a wash).

The syntax looks like this:

```

1 .macro macroname
2     instr1 a, b, c
3     instr2, b, d
4 # etc.
5 .end_macro
6
7 # or with parameters
8 .macro macroname(%arg1)
9     instr1    a, %arg1
10    instr2    c, d, e
11 # etc.
12 .end_macro

```

Some common examples are using them to print strings:

```
1 .macro print_str_label(%x)
2     li      $v0, 4
3     la      $a0, %x
4     syscall
5 .end_macro
6
7 .macro print_str(%str)
8 .data
9 str: .asciiz %str
10 .text
11     li      $v0, 4
12     la      $a0, str
13     syscall
14 .end_macro
15
16 .data
17
18 str1: .asciiz "Hello 1\n"
19
20 .text
21 # in use:
22     print_str_label(str1)
23
24     print_str("Hello World\n")
25
26     ...
```

You can see an example program in [macros.s](#).

Unfortunately, as far as I can tell, SPIM does not support function style macros despite what MARS's documentation implies about using a \$ instead of a % for arguments.

Switch-Case Statements

It is relatively common in programming to compare an integral type variable (ie basically any built-in type but float and double) against a bunch of different constants and do something different based on what it matches or if it matches none.

This could be done with a long **if-else-if** chain, but the longer the chain the more likely the programmer is to choose a **switch-case** statement instead.

Here's a pretty short/simple example in C:

```
1     printf("Enter your grade (capital): ");
2     int grade = getchar();
3     switch (grade) {
4     case 'A': puts("Excellent job!"); break;
```

```

5  case 'B': puts("Good job!"); break;
6  case 'C': puts("At least you passed?"); break;
7  case 'D': puts("Probably should have dropped it..."); break;
8  case 'F': puts("Did you even know you were signed up for the class?"); break;
9  default: puts("You entered an invalid grade!");
10 }

```

You could translate this to its equivalent `if-else` chain and handle it like we cover in the chapter on branching. However, imagine if this switch statement had a dozen cases, two dozen etc. The MIPS code for that quickly becomes long and ugly.

So what if we implemented it in MIPS the same way it is semantically in C? The same way compilers often (but not necessarily) use? Well, before we do that, what is a switch actually doing? It is *jumping* to a specific case label based on the value in the specified variable. It then starts executing, falling through any other labels, till it hits a `break` which will jump to the end of the switch block. If the value does not have its own case label, it will jump to the default label.

Compilers handle it by creating what's called a jump table, basically an array of label addresses, and using the variable to calculate an index in the table to use to jump to.

The C equivalent of that would look like this:

```

1  #include <stdio.h>
2
3
4  // This compiles with gcc, uses non-standard extension
5  // https://gcc.gnu.org/onlinedocs/gcc/Labels-as-Values.html
6
7  int main()
8  {
9
10     // jump table
11     void* switch_table[] =
12     { &a_label, &b_label, &c_label, &d_label, &default_label, &f_label };
13
14     printf("Enter your grade (capital): ");
15     int grade = getchar();
16     grade -= 'A'; // shift to 0
17
18     if (grade < 0 || grade > 'F' - 'A')
19         goto default_label;
20
21     goto *switch_table[grade];
22
23 a_label:
24     puts("Excellent job!");
25     goto end_switch;
26
27 b_label:
28     puts("Good job!");

```

```

29     goto end_switch;
30
31 c_label:
32     puts("At least you passed?");
33     goto end_switch;
34
35 d_label:
36     puts("Probably should have dropped it...");
37     goto end_switch;
38
39 f_label:
40     puts("Did you even know you were signed up for the class?");
41     goto end_switch;
42
43 default_label:
44     puts("You entered an invalid grade!");
45
46
47 end_switch:
48
49
50     return 0;
51 }

```

The `&&` and `goto *var` syntax are actually not standard C/C++ but are GNU extensions that are supported in gcc (naturally) and clang, possibly others.^[2]

First, notice how the size of the jump table is the value of the highest valued label minus the lowest + 1. That's why we subtract the lowest value to shift the range to start at 0 for the indexing. Second, any values without labels within that range are filled with the `default_label` address. Third, there has to be an initial check for values outside the range to jump to default otherwise you could get an error from an invalid access outside of the array's bounds.

The same program/code in MIPS would look like this:

```

1 .data
2
3 a_str: .asciiz "Excellent job!\n"
4 b_str: .asciiz "Good job!\n"
5 c_str: .asciiz "At least you passed?\n"
6 d_str: .asciiz "Probably should have dropped it...\n"
7 f_str: .asciiz "Did you even know you were signed up for the class?\n"
8
9 invalid_str: .asciiz "You entered an invalid grade!\n"
10
11 enter_grade: .asciiz "Enter your grade (capital): "
12
13 switch_labels: .word a_label, b_label, c_label, d_label, default_label, f_label
14
15 .text

```

```

16
17 main:
18
19     li    $v0, 4
20     la    $a0, enter_grade
21     syscall
22
23     li    $v0, 12    # read char
24     syscall
25
26     li    $t2, 5     # f is at index 5
27
28     la    $t0, switch_labels
29     addi   $t1, $v0, -65    # t1 = grade - 'A'
30     blt    $t1, $0, default_label    # if (grade-'A' < 0) goto default
31     bgt    $t1, $t2, default_label    # if (grade-'A' > 5) goto default
32
33     sll    $t1, $t1, 2    # offset *= 4 (sizeof(word))
34     add    $t0, $t0, $t1    # t0 = switch_labels + byte_offset =
    &switch_labels[grade-'A']
35     lw     $t0, 0($t0)    # load address from jump table
36     jr     $t0            # jump to address
37
38 a_label:
39     la    $a0, a_str
40     j     end_switch
41
42 b_label:
43     la    $a0, b_str
44     j     end_switch
45
46 c_label:
47     la    $a0, c_str
48     j     end_switch
49
50 d_label:
51     la    $a0, d_str
52     j     end_switch
53
54 f_label:
55     la    $a0, f_str
56     j     end_switch
57
58 default_label:
59     la    $a0, invalid_str
60
61
62 end_switch:
63     li    $v0, 4
64     syscall
65

```

```
66  li      $v0, 10  # exit
67  syscall
```

It's easy to forget that `jr` does not actually stand for "jump return" or "jump ra," though it's almost always used for that purpose. Instead, it stands for "jump register" and we can use it to do the equivalent of the computed goto statement in C.

This example probably wasn't worth making switch style, because the overhead and extra code of making the table and preparing to jump balanced out or even outweighed the savings of a branch instruction for every case. However, as the number of options increases, the favor tilts toward using a jump table like this as long as the range of values isn't too sparse. If the range of values is in the 100's or 1000's but you only have cases for a dozen or so, then obviously it isn't worth creating a table that large only to fill it almost entirely with the default label address.

To reiterate, remember it is not about the magnitude of the actual values you're looking for, only the difference between the highest and lowest because $\text{high} - \text{low} + 1$ is the size of your table.

Command Line Arguments

Command line arguments, also known as program arguments, or command line parameters, are strings that are passed to the program on startup. In high level languages like C, they are accessed through the parameters to the main function (naturally):

```
1  #include <stdio.h>
2
3  int main(int argc, char** argv)
4  {
5      printf("There are %d command line arguments:\n", argc);
6
7      for (int i=0; i<argc; i++) {
8          printf("%s\n", argv[i]);
9      }
10
11     return 0;
12 }
```

As you can see, `argc` contains the number of parameters and `argv` is an array of those arguments as C strings. If you run this program you'll get something like this:

```
$ ./args 3 random arguments
There are 4 command line arguments:
./args
3
random
arguments
```

Notice that the first argument is what you actually typed to invoke the program, so you always have

at least one argument.

MIPS works the same way. The number of arguments is in `$a0` and an array of strings is in `$a1` when main starts. So the same program in MIPS looks like this:

```
1 .data
2
3 there_are: .asciiz "There are "
4 arguments: .asciiz " command line arguments:\n"
5
6 .text
7 main:
8     move    $t0, $a0 # save argc
9
10    li      $v0, 4
11    la      $a0, there_are
12    syscall
13
14    move    $a0, $t0
15    li      $v0, 1 # print int
16    syscall
17
18    li      $v0, 4
19    la      $a0, arguments
20    syscall
21
22    li      $t1, 0 # i = 0
23    j       arg_loop_test
24
25 arg_loop:
26    li      $v0, 4
27    lw      $a0, 0($a1)
28    syscall
29
30    li      $v0, 11
31    li      $a0, 10 # '\n'
32    syscall
33
34    addi    $t1, $t1, 1 # i++
35    addi    $a1, $a1, 4 # argv++ ie a1 = &argv[i]
36 arg_loop_test:
37    blt     $t1, $t0, arg_loop # while (i < argc)
38
39    li      $v0, 10
40    syscall
```

This program works exactly like the C program when using SPIM:

```
$ spim -file args.s 3 random arguments
```

```
SPIM Version 8.0 of January 8, 2010
Copyright 1990-2010, James R. Larus.
All Rights Reserved.
See the file README for a full copyright notice.
Loaded: /usr/lib/spim/exceptions.s
There are 4 command line arguments:
args.s
3
random
arguments
```

Obviously the commands for SPIM itself are not included but the file name (args.s) takes the place as the "executable".

Unfortunately, MARS works differently, probably because it's more GUI focused. It does not pass the program/file name as the first argument, so you can actually get 0 arguments:

```
$ java -jar ~/Mars4_5.jar args.s pa 3 random arguments
MARS 4.5 Copyright 2003-2014 Pete Sanderson and Kenneth Vollmar

There are 3 command line arguments:
3
random
arguments
```

You can see that you have to pass "pa" (for "program arguments") to indicate that the following strings are arguments. In the GUI, there is an option in "Settings" called "Program arguments provided to MIPS program" which if selected will add a text box above the Text Segment for you to enter in the arguments to be passed.




Figure 1. Enable program arguments in MARS GUI

Edit

Execute

Text Segment

Program Arguments:

3 Random Arguments

Bkpt	Address	Code	Basic		
☐	0x00400000	0x00044021	addu \$8,\$0,\$4	11:	move \$t0, \$a0 # save argc
☐	0x00400004	0x24020004	addiu \$2,\$0,0x00000004	13:	li \$v0, 4
☐	0x00400008	0x3c011001	lui \$1,0x00001001	14:	la \$a0, there_are
☐	0x0040000c	0x34240000	ori \$4,\$1,0x00000000		
☐	0x00400010	0x0000000c	syscall	15:	syscall
☐	0x00400014	0x00082021	addu \$4,\$0,\$8	17:	move \$a0, \$t0
☐	0x00400018	0x24020001	addiu \$2,\$0,0x00000001	18:	li \$v0, 1 # print int
☐	0x0040001c	0x0000000c	syscall	19:	syscall
☐	0x00400020	0x24020004	addiu \$2,\$0,0x00000004	21:	li \$v0, 4
☐	0x00400024	0x3c011001	lui \$1,0x00001001	22:	la \$a0, arguments
☐	0x00400028	0x3424000b	ori \$4,\$1,0x0000000b		
☐	0x0040002c	0x0000000c	syscall	23:	syscall
☐	0x00400030	0x24090000	addiu \$9,\$0,0x00000000	25:	li \$t1, 0 # i = 0
☐	0x00400034	0x08100016	j 0x00400058	26:	j arg_loop_test
☐	0x00400038	0x24020004	addiu \$2,\$0,0x00000004	29:	li \$v0, 4

Data Segment

Address	Value (+0)	Value (+4)	Value (+8)	Value (+c)	Va
0x10010000	0x72656854	0x72612065	0x20002065	0x6d6d6f63	
0x10010020	0x0a3a7374	0x00000000	0x00000000	0x00000000	
0x10010040	0x00000000	0x00000000	0x00000000	0x00000000	
0x10010060	0x00000000	0x00000000	0x00000000	0x00000000	
0x10010080	0x00000000	0x00000000	0x00000000	0x00000000	
0x100100a0	0x00000000	0x00000000	0x00000000	0x00000000	
0x100100c0	0x00000000	0x00000000	0x00000000	0x00000000	
0x100100e0	0x00000000	0x00000000	0x00000000	0x00000000	
0x10010100	0x00000000	0x00000000	0x00000000	0x00000000	
0x10010120	0x00000000	0x00000000	0x00000000	0x00000000	
0x10010140	0x00000000	0x00000000	0x00000000	0x00000000	
0x10010160	0x00000000	0x00000000	0x00000000	0x00000000	
0x10010180	0x00000000	0x00000000	0x00000000	0x00000000	
0x100101a0	0x00000000	0x00000000	0x00000000	0x00000000	
0x100101c0	0x00000000	0x00000000	0x00000000	0x00000000	

←

→

0x10010000 (.data)

☒ Hexadecimal Addresses

Mars Messages

Run I/O

There are 3 command line arguments:

3

Random

Arguments

-- program is finished running --

Clear

Figure 2. Example using program arguments in MARS GUI

Delayed Branches and Delayed Loads

MIPS originally had a 1-instruction delay between when a branch or load was executed, and when it actually jumped or completed the load. By default, SPIM has both of these turned off but you can turn them on with the arguments `-delayed_branches` and `-delayed_loads` respectively, or use the `-bare` which turns on both, as well as disabling pseudoinstructions (ie it simulates the "bare" machine without any of the niceties). MARS is the same as SPIM in that it defaults to the more user friendly mode, but you can turn on delayed branches in the settings (but not delayed loads, that isn't supported as far as I can tell). From the command line you can turn it on with the `db` argument

and you can turn off pseudoinstructions with `np` (for the next section).

The only reason you'd probably ever want to do this to yourself is if your professor requires it for some unknown reason. There is some learning value to seeing how pseudoinstructions are turned into actual instructions, but there's absolutely nothing useful to be gained in turning on the delays.

To handle delayed branches, all you have to do is add a `nop` (No Operation) instruction after every branch or jump instruction (including `jal` and `jr`). The solution is identical for delayed loads, put a `nop` after every load instruction.

You could put other instructions there, instructions that you have above the jumps or loads that don't have to occur before. If you were doing a real project on real hardware and were trying to minimize program size and maximize speed, you would definitely do this (if that hardware had delayed branches/loads). However, it's a bad idea for two reasons. First, it takes longer to figure out which instructions you can move and why, and it makes your program harder to read since you naturally don't expect an instruction to occur before the previous one. Second, and more importantly, it means your program will no longer work if you turn delayed branches/loads off again. Using a `nop` means the change is quick, easy, and you can still run it in the simulator's default mode if needed.

One final note, and it relates to the following section. Even though `nop` is supposedly represented by all 0's, and MARS supports it even with pseudoinstructions disabled, SPIM does not support it in bare mode. This means you have to use some other instruction that has no side effects, like `or $0, $0, $0`, which will work as a `nop` under all conditions and simulators.

No Pseudoinstructions Allowed

Far more common than requiring delayed loads is forbidding pseudoinstructions, either all of them, or some subset of them. This forces us to explicitly write what actually happens when you use those pseudoinstructions.

Table 7. Pseudoinstruction Equivalents

Pseudoinstruction	Example Use	Equivalence
Load Immediate	<pre>li \$t0, 42</pre>	<pre>ori \$t0, \$0, 42 # or addi \$t0, \$0, 42</pre>
Move	<pre>move \$t0, \$t1</pre>	<pre>or \$t0, \$0, \$t1 # or add \$t0, \$0, \$t1</pre>
No Operation	<pre>nop</pre>	<pre># anything with \$0 as dest # reg will work, I prefer or \$0, \$0, \$0</pre>

Pseudoinstruction	Example Use	Equivalence
Load Address	<code>la \$t0, label</code>	<pre># use 0x1001 for MARS lui \$t0, 0x1000 ori \$t0, \$t0, byte_offset</pre>
Branch Less Than	<code>blt \$t0, \$t1, label</code>	<pre># t2 = t0 < t1 slt \$t2, \$t0, \$t1 bne \$t2, \$0, label</pre> <pre># or if you're counting up and # you know they'll be equal, # you can simplify to bne \$t0, \$t1, label</pre>
Branch Greater Than	<code>bgt \$t0, \$t1, label</code>	<pre># flip order to get > slt \$t2, \$t1, \$t0 bne \$t2, \$0, label</pre>
Branch Less Than or Equal	<code>ble \$t0, \$t1, label</code>	<pre># test for < and = separately slt \$t2, \$t0, \$t1 bne \$t2, \$0, label beq \$t0, \$t1, label</pre> <pre># or add 1 to change <= to < # use a spare reg if you need # to preserve the original value addi \$t3, \$t1, 1 slt \$t2, \$t0, \$t3 bne \$t2, \$0, label</pre>
Branch Greater Than or Equal	<code>bge \$t0, \$t1, label</code>	<pre># test for opposite # and branch on failure slt \$t2, \$t0, \$t1 beq \$t2, \$0, label</pre>

You can see how you use the non-pseudoinstructions to match the same behavior, and there's often (usually) more than one way. Of all of them, **ble** is the worst, because what was 1 instruction becomes 3, and you'll sometimes need an extra register to hold the "plus 1" value if you still need the original.

Another thing I should comment on is the **la** equivalence. The reason it is a pseudoinstruction in

the first place is that an address is 32 bits. That's also the size of a whole instruction. Clearly there's no way to represent a whole address and anything else at the same time. The lower left corner of the greensheet has the actual formats of the 3 different types of instructions and even the jump format still needs 6 bits for the opcode. This is why `lui` exists, in order to facilitate getting a full address into a register by doing it in two halves, 16 + 16. The lower 16 can be placed with `addi` or `ori` after the `lui`.

That begs the question, what actually goes in the upper half? Well, since we're dealing with addresses in the `.data` section, the upper portion should match the upper part of address of the `.data` section. In SPIM the `data` section starts at `0x10000000` in bare mode. In normal mode it is `0x10010000`, but you'd be able to use `la` in normal mode. However, in MARS it is always `0x10010000`, so you won't be able to have it work correctly in bare SPIM and MARS without changing that back and forth.

But what about the lower part of the address? This involves counting the bytes from the top of `.data` to the label you want. If all you have is words, halves, floats, doubles, or space (with a round number), that's fairly easy, but the second you have strings between the start and the label you want, it's a bit more painful. This is why I recommend putting any string declarations at the bottom so at least any other globals will have nice even offsets. Also, if you have a bunch of globals, it doesn't hurt to count once and put the offsets in comments above each label so you don't forget. Of course, none of this matters if you're allowed to just use `la` which is true the vast majority of the time.

Let's look at a small example. We'll convert the `args.s` from above (reproduced here for convenience) to bare mode:

```
1 .data
2
3 there_are: .asciiz "There are "
4 arguments: .asciiz " command line arguments:\n"
5
6 .text
7 main:
8     move    $t0, $a0 # save argc
9
10    li      $v0, 4
11    la      $a0, there_are
12    syscall
13
14    move    $a0, $t0
15    li      $v0, 1 # print int
16    syscall
17
18    li      $v0, 4
19    la      $a0, arguments
20    syscall
21
22    li      $t1, 0 # i = 0
23    j      arg_loop_test
```

```

24
25 arg_loop:
26     li      $v0, 4
27     lw      $a0, 0($a1)
28     syscall
29
30     li      $v0, 11
31     li      $a0, 10    # '\n'
32     syscall
33
34     addi     $t1, $t1, 1    # i++
35     addi     $a1, $a1, 4    # argv++ ie a1 = &argv[i]
36 arg_loop_test:
37     blt      $t1, $t0, arg_loop # while (i < argc)
38
39     li      $v0, 10
40     syscall

```

So we need to change the **moves**, **li**'s, **la**'s, and the **j**. after any branches or loads.

```

1 # in bare mode the .data section starts at 0x10000000
2 # without -bare, spim puts it at 0x10010000, so your lui will change
3 # but if you weren't doing bare you'd probably just be using la
4
5 # MARS .data always starts at 0x10010000, whether pseudoinstructions
6 # or delayed branching are on or not.
7
8 .data
9
10 there_are: .asciiz "There are "
11 arguments: .asciiz " command line arguments:\n"
12
13 # run with spim -bare -file args_bare.s
14
15 .text
16
17 .globl main
18 main:
19     or      $t0, $0, $a0 # save argc
20
21     ori     $v0, $0, 4
22     lui     $a0, 0x1000    # there_are is at beginning of data so just lui, lower
    is 0
23     syscall
24
25     or      $a0, $0, $t0
26     ori     $v0, $0, 1    # print int
27     syscall
28
29     ori     $v0, $0, 4

```

```

30    lui    $a0, 0x1000
31    ori    $a0, $a0, 11    # 11 is length in bytes of "There are " 10 chars + '\0'
32    #la    $a0, arguments
33    syscall
34
35    ori    $t1, $0, 0    # i = 0
36    j      arg_loop_test
37    or     $0, $0, $0
38
39 arg_loop:
40    ori    $v0, $0, 4    # print string for argv[i]
41    lw     $a0, 0($a1)
42    or     $0, $0, $0    # nop
43    syscall
44
45    ori    $v0, $0, 11
46    ori    $a0, $0, 10    # '\n'
47    syscall
48
49    addi   $t1, $t1, 1    # i++
50    addi   $a1, $a1, 4    # argv++ ie a1 = &argv[i]
51 arg_loop_test:
52    bne    $t1, $t0, arg_loop    # while (i != argc)
53    or     $0, $0, $0
54
55    ori    $v0, $0, 10
56    syscall

```

Following the table, you can see the **mv** became **or**, the **li** became **ori**, we added **nop** instructions in the form of **or** with **\$0** as the destination, the **blt** became **bne**, and lastly the **la** became **lui** plus an **ori** if necessary for the byte offset. The **blt** to **bne** is worth noting, because there's no reason to do a complicated transformation if you don't have to and **bne** accomplishes the same thing here.

Exercises

You can support the book and purchase the chapter exercise solutions from [my store](#), [Leanpub](#), or [itch.io](#).

1. Convert the exercises from chapter 5 to run in **spim -bare** mode and/or in MARS with delayed branches and no pseudoinstructions (**java -jar ~/Mars_4.5.jar db np file.s** on the command line).
2. Convert the following C code to MIPS using a jump table (Note in C/C++ enum values start at 0 and go up by one unless the user manually assigns a value, in which case it continues counting up from there).

```

1 enum { STATE0, STATE1, STATE2, STATE3, STATE14 = 14, STATE42 = 42, STATE43,
    STATE44 };
2

```

```

3
4 int main()
5 {
6     int num;
7     do {
8         printf("Enter a number between 0 and 50: ");
9         scanf("%d", &num);
10    } while (num < 0 || num > 50);
11
12    switch (num) {
13    case STATE0:
14        puts("Zilch");
15        break;
16    case STATE1:
17        puts("Uno");
18        break;
19    case STATE2:
20        puts("Dos");
21        break;
22    case STATE3:
23        puts("Tres");
24        break;
25    case STATE14:
26        puts("Catorce");
27    case STATE42:
28        puts("The answer to life, the universe, and everything.");
29    case STATE43:
30        puts("Off by one");
31    case STATE44:
32        puts("4 * 11?");
33        break;
34    }
35
36    puts("Thanks for playing!");
37
38
39    return 0;
40 }

```

[1] When I find the post I read years ago about how using tabs saves CO2 I'll put it here, but I'm joking. I use tabs because it makes sense and there are accessibility reasons too: https://www.reddit.com/r/javascript/comments/c8drjo/nobody_talks_about_the_real_reason_to_use_tabs/

[2] <https://gcc.gnu.org/onlinedocs/gcc/Labels-as-Values.html>

References and Useful Links

- [Greensheet](#)
- [MARS syscall list](#)
- [learnxinyminutes page](#)
- [Randolph-Macon College MIPS Reference](#)
- [CCSU MIPS Reference](#)

Supporters

Corporate

MIPS Assembly Programming	Robert Winkler	supporter3
---	--------------------------------	------------

Platinum

supporter4	supporter5	PortableGL
supporter7		

Gold

sdl_img	Nick, Columbia CS Student	Dr. Michael Rogers Assistant Professor Computer Science University of Wisconsin Oshkosh
supporter11	supporter12	supporter13

Silver

Atlas Gong https://cs.mcgill.ca/~agong2	supporter15	supporter16
supporter17	supporter18	supporter19
supporter20		

Bronze

CCSF cs student 0x01	supporter22	supporter23
supporter24	supporter25	supporter26
supporter27	supporter28	supporter29
supporter30		